

ELENA REVICHEVA

Applied AI Engineer

Production LLM systems for CRM, approved outreach, and AI-search visibility

aipa@aideazz.xyz · +507 616 66 716 · LinkedIn · GitHub · aideazz.xyz/portfolio

Contract · 20–30 hours/week · Panama (UTC-5) · daily overlap 09:00–13:00 ET

SUMMARY

I design, deploy, and stay on call for the automation that runs a business when nobody is watching: find the customer, draft the reply, write the CRM, report what happened. Twelve production systems on one Oracle Cloud VM for eighteen months. Before that, seven years as Deputy CEO running regulated digital-transformation programs at board level, so I can hold the same conversation with the engineer and the person signing.

SELECTED WORK

AI Visibility Audit API — public product

aideazz.xyz/api

Live scoring for how discoverable a site is to ChatGPT, Perplexity and Claude, not only Google. Same GEO/AEO engine that scores my own hub 100/100. Production floor from Oracle logs, rounded down: 420+ audits, 14,000+ signals, 210+ sites, median 85. Each check returns the fix and why it matters (34 why / 4 fix on a live stripe.com run).

Approved-send growth loop — HubSpot + Resend

The acquisition loop I run in production: research, qualify, draft, one human tap, send, then delivery and open written back to the CRM. 1,900+ deals, 989 contacts, 1,411 companies attributed by the agents. Send is fail-closed: if a PDF or signed .docx fails to load, the email is refused so it cannot claim an attachment it did not carry.

EspaLuz — paying WhatsApp / Telegram tutor

Bilingual relocation and language tutor on WhatsApp and Telegram. Persistent per-student memory (RAG + pgvector). PayPal subscriptions. Early users in 19 countries. I am the architect and the on-call.

Fail-closed publishing and a five-provider LLM chain

aideazz.xyz/ai-ops-wiki.html

Every model call in the fleet has an ordered fallback (Anthropic, OpenAI, Gemini, Grok, Groq). When Groq deprecated the models I used, the fleet kept serving — a config change, not an outage. The daily publisher will not print a number it cannot trace; if the model decorated the facts, the day stayed silent. A 130-test eval harness (unit, integration, golden-set) runs in under a minute at \$0 API cost.

HOW I WORK

I am the architect, the reviewer and the on-call. I judge a new data source by running it through the live gate before integrating it: the source added in August cleared 76% against ~21% for the rest of the fleet. I write the failure down and name it — clients can read how I debug before they hire me.

AVAILABLE FOR

- Production LLM systems with fallback, evals and on-call ownership
- CRM and approved-send outreach (HubSpot, Resend, human-in-the-loop)
- AI-search visibility (GEO/AEO) audits and engines — ChatGPT, Perplexity, Claude
- Fail-closed publishing and internal AI operations for a founder-led team

STACK

Python · TypeScript · Node.js · FastAPI · LangGraph · RAG (pgvector) · HubSpot · Make · Resend · Playwright · PostgreSQL · Docker · Oracle Cloud (OCI)

EARLIER

Deputy CEO & Chief Legal Officer — JSC “E-GOV OPERATOR” · 2011–2018

Seven years at board level on large-scale public-sector digital transformation: IT, legal and compliance in a heavily regulated environment.

EDUCATION

MA Social Psychology — Penza State University · 2018. Presidential Program for Executive Management — RANEPa · 2015. Internship — Nyskapingsparken Innovation Park, Bergen, Norway.

LANGUAGES

English — fluent (professional) · Spanish — intermediate · Russian — native · French — elementary