

ELENA REVICHEVA

Technical Product Owner — AI & Agentic Platforms | Forward-Deployed | LATAM

Panama City, Panama (UTC-5) — already in LATAM, no relocation needed · aipa@aideazz.xyz · +507 616 66 716

aideazz.xyz/portfolio · linkedin.com/in/elenarevicheva · github.com/ElenaRevicheva

PROFILE

Seven years as a **Deputy CEO** owning large-scale digital transformation at board level, then two years building and operating **AI products end to end** — I write the roadmap, ship it, carry the pager, and answer to the numbers. Twelve systems run in production today on one cloud VM: I am the product owner, the reviewer and the on-call. What I do daily is exactly the Hilbert shape — **data pipelines and ML infrastructure with an agent on top**, sold to non-technical operators who buy an outcome, not a model. I can hold the architecture conversation with engineers and then explain the trade-off to the person signing the cheque, in the same meeting.

PRODUCT OWNERSHIP, 0 → 1 (each shipped, live, and measured)

- **Autonomous job-and-lead pipeline (0→1, live 18 months)**. Owned the whole surface: sourcing, scoring rubric, the LLM judge, and the CRM it writes into. **1,898 records** processed. Defined the success metric myself, then acted on it — when outcome data showed one role category produced **58% of rejections but only 17% of positives**, I cut it from the targeting in a day and re-pointed the paid search queries. That is the roadmap coming from measured customer reality, not opinion.
- **AI Visibility Audit API (0→1, public product)**. Scoped and shipped a scoring product: **34 automated checks in 4 weighted categories**, returning a grade and a prioritised fix list. Designed the weighting — the judgement call of what counts and how much — and the output non-technical buyers act on. aideazz.xyz/api
- **Bilingual AI tutor with paying subscribers (0→1)**. Built from my own family's need as an expat in Panama, through to billing. Persistent per-user memory, voice, OCR, subscription payments. Early revenue, honestly small — included because I owned it from problem to paid product.
- **Growth & discovery engine (0→1)**. Daily automated publishing, structured data, attribution from first touch to CRM record, and inbound triage with urgency scoring. Instrumented end to end, so every claim about it is a number I can show rather than a story.

JUDGEMENT ON AI PRODUCTS — evals, non-determinism, ship/no-ship

- **Evals as a product gate, not a vanity metric**. **529 passing tests** across unit, integration and golden-set layers, running in about a minute at \$0 API cost. It is how I answer "is the model good enough to ship" with evidence instead of a feeling.
- **Non-determinism handled in the architecture**. Every LLM call runs a **five-provider fallback chain** ordered per use case across six products. When a provider deprecated the model I depended on, the fleet kept serving and the migration was a config change, not an outage.
- **Abstention over hallucination**. Research features are built to return "not enough public information" rather than invent — in the last production run, 2 of 18 correctly refused. In a product customers make revenue decisions on, a confident wrong answer costs more than a blank one.
- **Public postmortems**. **21 published incident write-ups** at aideazz.xyz/ai-ops-wiki.html — you can read how I reason about failure before you interview me.

FORWARD-DEPLOYED & COMMERCIAL

I run my own discovery and onboarding: direct outreach, qualification calls and follow-up through a CRM I built the automation for — so I hear the problem behind the request and know which version of it someone will pay for. Seven years at board level in a heavily regulated environment means I am comfortable being the person who says what the trade-off is, in front of people who did not want to hear it. Based in Panama, working across US and European time zones daily.

TECHNICAL FOUNDATION (enough to scope ruthlessly and prototype)

Data: SQL, PostgreSQL, Oracle Autonomous Database, pgvector, ingestion pipelines, RAG retrieval · **AI:** Claude, OpenAI, Gemini, Grok, Groq; multi-provider routing; agent orchestration; prompt systems; eval harnesses · **Build:** Python, TypeScript, FastAPI, Node.js, REST/webhooks · **Ops:** OCI, PM2, systemd, cron, monitoring and alerting · **Analytics:** GA4, UTM attribution, Search Console, HubSpot · **Method:** AI-assisted development (Claude Code, Cursor) as the daily workflow — the reason one person ships at this rate.

EXECUTIVE BACKGROUND

Deputy CEO & Chief Legal Officer — JSC "E-GOV OPERATOR" | 2011–2018

Led public-sector digital transformation at board level for seven years; ran cross-functional IT, legal and compliance teams in a heavily regulated environment. Owner's mindset is not a phrase here — I have made calls without perfect information and carried the consequences.

Deputy CEO, Business Development — Fundery LLC (fintech) | 2017–2018

EDUCATION & LANGUAGES

Presidential Program for Executive Management, RANEPa Moscow (2015) · MA Social Psychology, Penza State University (2018) · Polkadot Blockchain Academy, PBA-X Wave 3 (2025)

Languages: Russian (native) · English (fluent) · **Spanish (working — resident in Panama)** · French (elementary)