

ELENA REVICHEVA

Technical Product Manager

AI products where correctness and trust come first

aipa@aideazz.xyz · +507 616 66 716 · LinkedIn · GitHub · aideazz.xyz/portfolio

Full-time or contract · Remote worldwide or on-site in Panama City · UTC-5

SUMMARY

I design intelligent systems and the experience around them — from the agents that run a business to the films they make.

Seven years at board level as Deputy CEO and Chief Legal Officer of a public-sector e-government operator — IT, legal and compliance in a heavily regulated environment — and Deputy CEO for business development at a fintech. Since May 2025 I have owned AI products end to end: roadmap, acceptance criteria, trade-offs, launch and the on-call. My rule for AI is that it earns its place: a human approves every outbound send, no number is published unless it traces back to a record, and a model's decisions are measured against real human decisions before I rely on them.

EXPERIENCE

Founder & AI Product and Solutions Lead — Aideazz.xyz · Panama / Remote · 2025–Present

Founder-led AI lab. I design, build and operate its production AI systems end to end — the architect, the reviewer and the on-call.

AI screening judge — measured on real decisions before it ships

aideazz.xyz/ai-ops-wiki.html

An LLM judge screens the jobs my search agent finds; what the model may not decide is enforced in code — a location or pay it misreads is overruled from the posting itself. Each of my own decisions is stored with the posting it was made on, and the judge is measured against them: 14 of 20 rejections and 8 of 18 applications agree. A retrieval (RAG) upgrade scored lower on that test, so it ships switched off.

Fail-closed publishing and a five-provider LLM chain

aideazz.xyz/ai-ops-wiki.html

Every model call in the fleet has an ordered fallback (Anthropic, OpenAI, Gemini, Grok, Groq). When Groq deprecated the models I used, the fleet kept serving — a config change, not an outage. The daily publisher will not print a number it cannot trace; if the model decorated the facts, the day stayed silent. A 130-test eval harness (unit, integration, golden-set) runs in under a minute at \$0 API cost.

Approved-send growth loop — HubSpot + Resend

The acquisition loop I run in production: research, qualify, draft, one human tap, send, then delivery and open written back to the CRM. 1,900+ deals, 989 contacts, 1,411 companies attributed by the agents. Send is fail-closed: if a PDF or signed .docx fails to load, the email is refused so it cannot claim an attachment it did not carry.

Operational Co-Founder — OmniBazaar, decentralised e-commerce · 2024–2025

PROOF HUB

Portfolio	aideazz.xyz/portfolio	Atlas	atlas.html
Wiki	ai-ops-wiki.html	Ops SOP	sop-ai-ops.html
Blog	aideazz.xyz/blog	Podcast	podcast.aideazz.xyz
EspaLuz WA	wa.me/50766623757	Atuona Studio	atuona.xyz/aifilmstudio

HOW I WORK

I operate an AI-native development environment where specialized agents handle much of the implementation execution. I own requirements, architecture, orchestration, evaluation, deployment, monitoring and production decisions. I judge a new data source by running it through the live gate before integrating it: the source added in August cleared 76% against ~21% for the rest of the fleet. I write the failure down and name it — clients can read how I debug before they hire me.

AVAILABLE FOR

- AI product ownership in regulated, high-consequence domains
- LLM features framed as evals: offline replay on real decisions, human review, fail-closed defaults

- Design reviews with engineering: failure modes, retries, idempotent writes, read-back audits
- Compliance-aware product decisions (Chief Legal Officer background, blockchain regulation)

STACK

Python · TypeScript · Node.js · FastAPI · LangGraph · RAG (pgvector) · HubSpot · Make · Resend · Playwright · PostgreSQL · Docker · Oracle Cloud (OCI)

EARLIER

Deputy CEO & Chief Legal Officer — JSC “E-GOV OPERATOR” · 2011–2018

Seven years at board level on large-scale public-sector digital transformation: IT, legal and compliance in a heavily regulated environment.

Deputy CEO, Business Development — Fundery LLC (fintech) · 2017–2018

EDUCATION

MA Social Psychology — Penza State University · 2018. Blockchain Regulation — MGIMO, Moscow · 2017. Presidential Program for Executive Management — RANEPa · 2015. Polkadot Blockchain Academy, PBA-X Wave 3 · 2025. How-To-DAO Cohort · 2025. Anthropic Academy — Claude Certification Program · 2026, in progress. Internship — Nyskapingsparken Innovation Park, Bergen, Norway.

LANGUAGES

English — fluent (professional) · Spanish — intermediate · Russian — native · French — elementary