

ELENA REVICHEVA

AI-Native Operations

Governance · standards · operating cadence · evaluation · agentic execution

Panama City, Panama · UTC-5 · Remote worldwide · aipa@aideazz.xyz · +507 616 66 716 · linkedin.com/in/elenarevicheva · github.com/ElenaRevicheva · aideazz.xyz/portfolio

SUMMARY

Seven years as Deputy CEO for Corporate Governance and Chief Legal Officer of a regional e-government operator: I arranged the board meetings, drove decisions through to implementation and checked the services we moved online against the law. Today I design the standards, controls and evaluations behind production agentic systems, with a human approval on every consequential action.

I operate an AI-native development environment where specialized agents handle much of the implementation execution. I own requirements, architecture, orchestration, evaluation, deployment, monitoring and production decisions.

EXPERIENCE

Founder & AI Product and Solutions Lead — Aideazz AI Lab · Panama / Remote · 2025–Present

Founder-led AI lab. I direct the agents that build and run its production systems, and own review, release and on-call.

Standards, decisions and cadence

- Wrote the coordination protocol three AI coding agents follow even though they cannot see each other's chats: claim before touching shared systems, a handoff block on every pause (done / next / verified by / risk), every work product in a committed file or the CRM — never only in a chat — and a start-of-session checklist built from failures that shipped.
- Each CRM stage is named for whose move it is; one approved send moves the deal, stamps delivery and books the follow-up itself. A read-back audit checks every priority deal against a six-part completeness standard (7 of 7 complete today).
- Built an automated morning page that replaced opening CRM records one by one: the day's priority queue, each with a cited company brief where sources exist.
- I publish our own failures: 20+ named postmortems in a public [AI Operations Wiki](#), plus an [AI Operations SOP](#).

Production operations

- 15 long-running production services and 20 scheduled jobs on one cloud VM, all auto-restarting, with a 5-provider LLM fallback chain so a provider failure degrades service instead of stopping it.
- Approved-send loop on HubSpot + Resend: research → qualify → draft → one human approval → send → delivery and opens written back to the CRM. Fail-closed: if a listed PDF or .docx contract cannot load, the email is refused rather than claiming an attachment it did not carry.
- I audit from the receiving system, not the sender's log. That exposed EspaLuz's trial-signup → CRM integration (our language-tutor bot) silently broken since launch; root-caused, fixed and backfilled all 28 missing users.

Evaluation and metrics

- I challenge the metric, not just the model: the first replay of my job-screening judge was grading it against blank postings, so its score measured the blank. Since September each decision is stored with its real posting and replayed as an offline evaluation set before a judge change ships.
- A proposed RAG upgrade did not beat the baseline on a 38-decision replay, so it stayed off. Location and pay are hard gates in code that run before the model, so the model cannot wave through a job that fails them.
- 600+ automated tests on the job-discovery pipeline (unit, integration, golden-set), running in about a minute; 11,000+ job postings qualified by the agentic pipeline.

AI Visibility Audit API — public product · aideazz.xyz/api

- I designed the rubric: 34 automated checks in 4 weighted categories, scored to a grade plus a prioritised fix plan. Every check explains why it matters; failed checks return the fix. 420+ assessments across 210+ sites (production floor from logs).

Deputy CEO for Corporate Governance & Chief Legal Officer — JSC “E-GOV OPERATOR” · Penza region · 2011–2018

A ~60-person company moving the delivery of public services into electronic form for the regional government.

- Owned the board process: arranged board meetings for stakeholders drawn mainly from the Penza regional government and drove decisions through to implementation.
- Owned the legal standard for the programme: made sure the public services the company built in electronic form were eligible under existing law, and managed those projects' legal workstream.
- Conducted due diligence on public contracts and handled the company's court litigation.

Operational Co-Founder — OmniBazaar, decentralized e-commerce · 2024–2025

Deputy CEO, Business Development — Fundery LLC, fintech · 2017–2018

Figures are production floors from logs and the live CRM, counted September 2026.

TECHNICAL FLUENCY

AI & orchestration: LangGraph · RAG (pgvector) · multi-provider LLM routing · evaluation systems · human-in-the-loop workflows ·
Engineering: Python · TypeScript · Node.js · FastAPI · PostgreSQL · Docker · Playwright · **Operations:** HubSpot · Resend · Make · Oracle
Cloud (OCI) · Git/GitHub · APIs · **AI-native development:** Claude Code · Cursor · specialized coding and operational agents

EDUCATION & PROFESSIONAL DEVELOPMENT

MA, Social Psychology — Penza State University · 2018 · Presidential Program for Executive Management — RANEPa · 2015 · Blockchain
Regulation — MGIMO, Moscow · 2017 · Polkadot Blockchain Academy, PBA-X Wave 3 · 2025 · How-To-DAO Cohort · 2025 · Anthropic
Academy — Claude Certification Program · 2026, in progress · Internship — Nyskapingsparken Innovation Park, Bergen, Norway

LANGUAGES

English — fluent (professional) · Spanish — intermediate (working; resident in Panama) · Russian — native · French — elementary

LINKS

aideazz.xyz/portfolio · [AI Operations Wiki](#) · [AI Operations SOP](#) · [AI Visibility Audit](#) · [GitHub](#) · [LinkedIn](#)