

ELENA REVICHEVA

AI Solutions Architect — production systems, agentic automation, integrations

Panama City, Panama (UTC-5) · aipa@aideazz.xyz · +507 616 66 716

aideazz.xyz/portfolio · linkedin.com/in/elenarevicheva · github.com/ElenaRevicheva

PROFILE

I design AI and automation solutions and then stay on call for them. Seven years at board level before this, so I can scope against a business outcome rather than a wishlist, and say no to the parts that will not pay for themselves. Everything I describe below is running in production, not a prototype.

SELECTED WORK — shipped, live, and measured

- **Owned an autonomous job-and-lead pipeline 0 to 1, live since November 2025.** Sourcing, the scoring rubric, the LLM judge, and the CRM it writes into. **1,898 records** processed. I set the success metric and then acted on it: when outcome data showed one role category produced **58% of rejections but only 17% of positives**, I cut it from targeting in a day and re-pointed the paid search queries. Roadmap from measured reality, not opinion.
- **Shipped a public scoring API 0 to 1. 34 automated checks in 4 weighted categories**, returning a grade and a prioritised fix list. I designed the weighting — the judgement call of what counts and how much — and the output a non-technical buyer can act on.
- **Evals as a ship gate, not a vanity metric. 529 passing tests** across unit, integration and golden-set layers, running in about a minute at \$0 API cost. It is how I answer "is the model good enough to ship" with evidence rather than a feeling.
- **Non-determinism handled in the architecture.** Every LLM call runs a **five-provider fallback chain** ordered per use case across six products. When a provider deprecated the model I depended on, the fleet kept serving and the migration was a config change, not an outage.
- **I run the fleet, not just the code. 15 long-running services and 20 scheduled jobs** on a single cloud VM at \$0/month infrastructure, 24/7. I am the architect, the reviewer and the on-call — three-layer recovery so failures heal without me.
- **Forward-deployed by default.** I run my own discovery, onboarding and follow-up through a CRM I built the automation for, so I hear the problem behind the request and know which version of it someone will pay for.
- **I publish my own postmortems. 21 incident write-ups** in a public AI Ops wiki — you can read how I reason about failure before you interview me.

TECHNICAL FOUNDATION

Data: SQL, PostgreSQL, Oracle Autonomous Database, pgvector, ingestion pipelines, RAG retrieval · **AI:** Claude, OpenAI, Gemini, Grok, Groq; multi-provider routing; agent orchestration; prompt systems; eval harnesses · **Build:** Python, TypeScript, FastAPI, Node.js, REST/webhooks · **Ops:** OCI, PM2, systemd, cron, monitoring and alerting · **Analytics:** GA4, UTM attribution, Search Console, HubSpot · **Method:** AI-assisted development (Claude Code, Cursor) as the daily workflow.

EXECUTIVE BACKGROUND

Deputy CEO & Chief Legal Officer — JSC "E-GOV OPERATOR" | 2011–2018

Led public-sector digital transformation at board level for seven years; ran cross-functional IT, legal and compliance teams in a heavily regulated environment.

Deputy CEO, Business Development — Fundery LLC (fintech) | 2017–2018

EDUCATION & LANGUAGES

Presidential Program for Executive Management, RANEPa Moscow (2015) · MA Social Psychology, Penza State University (2018) ·

Polkadot Blockchain Academy, PBA-X Wave 3 (2025)

Languages: Russian (native) · English (fluent) · **Spanish (working — resident in Panama)** · French (elementary)