

ELENA REVICHEVA

AI Evaluation — evals, red-teaming, and ship/no-ship judgement

Panama City, Panama (UTC-5) · aipa@aideazz.xyz · +507 616 66 716

aideazz.xyz/portfolio · linkedin.com/in/elenarevicheva · github.com/ElenaRevicheva

PROFILE

I evaluate AI systems for a living because I operate them: I write the harnesses, set the gates, and decide when a model is good enough to ship. Seven years at board level before this means I can also explain why a failure matters commercially, not just that a test went red.

SELECTED WORK — shipped, live, and measured

- **Owned an autonomous job-and-lead pipeline 0 to 1, live since November 2025.** Sourcing, the scoring rubric, the LLM judge, and the CRM it writes into. **1,898 records** processed. I set the success metric and then acted on it: when outcome data showed one role category produced **58% of rejections but only 17% of positives**, I cut it from targeting in a day and re-pointed the paid search queries. Roadmap from measured reality, not opinion.
- **Eval as a ship gate, not a vanity metric.** **529 passing tests** across unit, integration and golden-set layers, running in about a minute at \$0 API cost. It is how I answer "is the model good enough to ship" with evidence rather than a feeling.
- **Non-determinism handled in the architecture.** Every LLM call runs a **five-provider fallback chain** ordered per use case across six products. When a provider deprecated the model I depended on, the fleet kept serving and the migration was a config change, not an outage.
- **Abstention designed in, over hallucination.** Research features return "not enough public information" rather than invent — in the last production run, **2 of 18** correctly refused. Where customers make decisions on the output, a confident wrong answer costs more than a blank one.
- **I publish my own postmortems.** **21 incident write-ups** in a public AI Ops wiki — you can read how I reason about failure before you interview me.

TECHNICAL FOUNDATION

Data: SQL, PostgreSQL, Oracle Autonomous Database, pgvector, ingestion pipelines, RAG retrieval · **AI:** Claude, OpenAI, Gemini, Grok, Groq; multi-provider routing; agent orchestration; prompt systems; eval harnesses · **Build:** Python, TypeScript, FastAPI, Node.js, REST/webhooks · **Ops:** OCI, PM2, systemd, cron, monitoring and alerting · **Analytics:** GA4, UTM attribution, Search Console, HubSpot · **Method:** AI-assisted development (Claude Code, Cursor) as the daily workflow.

EXECUTIVE BACKGROUND

Deputy CEO & Chief Legal Officer — JSC "E-GOV OPERATOR" | 2011–2018

Led public-sector digital transformation at board level for seven years; ran cross-functional IT, legal and compliance teams in a heavily regulated environment.

Deputy CEO, Business Development — Fundery LLC (fintech) | 2017–2018

EDUCATION & LANGUAGES

Presidential Program for Executive Management, RANEPA Moscow (2015) · MA Social Psychology, Penza State University (2018) · Polkadot Blockchain Academy, PBA-X Wave 3 (2025)

Languages: Russian (native) · English (fluent) · **Spanish (working — resident in Panama)** · French (elementary)