

ELENA REVICHEVA

AI Leadership — strategy, adoption and the systems underneath it

Panama City, Panama (UTC-5) · aipa@aideazz.xyz · +507 616 66 716

aideazz.xyz/portfolio · linkedin.com/in/elenarevicheva · github.com/ElenaRevicheva

PROFILE

Seven years as **Deputy CEO** leading digital transformation at board level across IT, legal and compliance in a heavily regulated environment — then, since May 2025, building the AI systems myself rather than delegating them. I can set the AI strategy and also tell you honestly which parts of it will fail, because I have operated them.

SELECTED WORK — shipped, live, and measured

- **Owned an autonomous job-and-lead pipeline 0 to 1, live since November 2025.** Sourcing, the scoring rubric, the LLM judge, and the CRM it writes into. **1,898 records** processed. I set the success metric and then acted on it: when outcome data showed one role category produced **58% of rejections but only 17% of positives**, I cut it from targeting in a day and re-pointed the paid search queries. Roadmap from measured reality, not opinion.
- **Non-determinism handled in the architecture.** Every LLM call runs a **five-provider fallback chain** ordered per use case across six products. When a provider deprecated the model I depended on, the fleet kept serving and the migration was a config change, not an outage.
- **Abstention designed in, over hallucination.** Research features return "not enough public information" rather than invent — in the last production run, **2 of 18** correctly refused. Where customers make decisions on the output, a confident wrong answer costs more than a blank one.
- **I run the fleet, not just the code.** **15 long-running services and 20 scheduled jobs** on a single cloud VM at \$0/month infrastructure, 24/7. I am the architect, the reviewer and the on-call — three-layer recovery so failures heal without me.
- **Forward-deployed by default.** I run my own discovery, onboarding and follow-up through a CRM I built the automation for, so I hear the problem behind the request and know which version of it someone will pay for.
- **Took a product from a personal problem to paying subscribers.** A bilingual AI tutor built from my own family's need as an expat in Panama, through to billing — persistent per-user memory, voice, OCR, subscriptions. Early revenue, honestly small; included because I owned it from problem to paid.
- **I publish my own postmortems.** **21 incident write-ups** in a public AI Ops wiki — you can read how I reason about failure before you interview me.

TECHNICAL FOUNDATION

Data: SQL, PostgreSQL, Oracle Autonomous Database, pgvector, ingestion pipelines, RAG retrieval · **AI:** Claude, OpenAI, Gemini, Grok, Groq; multi-provider routing; agent orchestration; prompt systems; eval harnesses · **Build:** Python, TypeScript, FastAPI, Node.js, REST/webhooks · **Ops:** OCI, PM2, systemd, cron, monitoring and alerting · **Analytics:** GA4, UTM attribution, Search Console, HubSpot · **Method:** AI-assisted development (Claude Code, Cursor) as the daily workflow.

EXECUTIVE BACKGROUND

Deputy CEO & Chief Legal Officer — JSC "E-GOV OPERATOR" | 2011–2018

Led public-sector digital transformation at board level for seven years; ran cross-functional IT, legal and compliance teams in a heavily regulated environment.

Deputy CEO, Business Development — Fundery LLC (fintech) | 2017–2018

EDUCATION & LANGUAGES

Presidential Program for Executive Management, RANEPa Moscow (2015) · MA Social Psychology, Penza State University (2018) · Polkadot Blockchain Academy, PBA-X Wave 3 (2025)

Languages: Russian (native) · English (fluent) · **Spanish (working — resident in Panama)** · French (elementary)