

ELENA REVICHEVA

Technical Product Owner — AI & Agentic Platforms | Forward-Deployed

Panama City, Panama (UTC-5) · aipa@aideazz.xyz · +507 616 66 716

aideazz.xyz/portfolio · linkedin.com/in/elenarevicheva · github.com/ElenaRevicheva

PROFILE

Seven years as a **Deputy CEO** owning digital transformation at board level, then two years building and operating **AI products end to end** — I write the roadmap, ship it, carry the pager, and answer to the numbers. What I do daily is data pipelines and ML infrastructure with an agent on top, sold to operators who buy an outcome, not a model. I hold the architecture conversation with engineers and then explain the trade-off to the person signing the cheque, in the same meeting.

SELECTED WORK — shipped, live, and measured

- **Owned an autonomous job-and-lead pipeline 0 to 1, live 18 months.** Sourcing, the scoring rubric, the LLM judge, and the CRM it writes into. **1,898 records** processed. I set the success metric and then acted on it: when outcome data showed one role category produced **58% of rejections but only 17% of positives**, I cut it from targeting in a day and re-pointed the paid search queries. Roadmap from measured reality, not opinion.
- **Shipped a public scoring API 0 to 1. 34 automated checks in 4 weighted categories**, returning a grade and a prioritised fix list. I designed the weighting — the judgement call of what counts and how much — and the output a non-technical buyer can act on.
- **Evals as a ship gate, not a vanity metric. 529 passing tests** across unit, integration and golden-set layers, running in about a minute at \$0 API cost. It is how I answer "is the model good enough to ship" with evidence rather than a feeling.
- **Abstention designed in, over hallucination.** Research features return "not enough public information" rather than invent — in the last production run, **2 of 18** correctly refused. Where customers make decisions on the output, a confident wrong answer costs more than a blank one.
- **Forward-deployed by default.** I run my own discovery, onboarding and follow-up through a CRM I built the automation for, so I hear the problem behind the request and know which version of it someone will pay for.
- **Built the discovery engine that makes a business quotable by AI assistants.** Structured data, AI-crawler access, answer-ready content and daily automated publishing — measured with a weekly citation probe across Google AI Overviews, Gemini, OpenAI search and Perplexity, plus first-touch-to-CRM attribution.
- **Took a product from a personal problem to paying subscribers.** A bilingual AI tutor built from my own family's need as an expat in Panama, through to billing — persistent per-user memory, voice, OCR, subscriptions. Early revenue, honestly small; included because I owned it from problem to paid.

TECHNICAL FOUNDATION

Data: SQL, PostgreSQL, Oracle Autonomous Database, pgvector, ingestion pipelines, RAG retrieval · **AI:** Claude, OpenAI, Gemini, Grok, Groq; multi-provider routing; agent orchestration; prompt systems; eval harnesses · **Build:** Python, TypeScript, FastAPI, Node.js, REST/webhooks · **Ops:** OCI, PM2, systemd, cron, monitoring and alerting · **Analytics:** GA4, UTM attribution, Search Console, HubSpot · **Method:** AI-assisted development (Claude Code, Cursor) as the daily workflow.

EXECUTIVE BACKGROUND

Deputy CEO & Chief Legal Officer — JSC "E-GOV OPERATOR" | 2011–2018

Led public-sector digital transformation at board level for seven years; ran cross-functional IT, legal and compliance teams in a heavily regulated environment.

Deputy CEO, Business Development — Fundery LLC (fintech) | 2017–2018

EDUCATION & LANGUAGES

Presidential Program for Executive Management, RANEPA Moscow (2015) · MA Social Psychology, Penza State University (2018) ·

Polkadot Blockchain Academy, PBA-X Wave 3 (2025)

Languages: Russian (native) · English (fluent) · **Spanish (working — resident in Panama)** · French (elementary)