

Eight private production repositories. I keep the Lab.

\$89,481

HUD point estimateNon-exclusive · products stay up · I retain copyright

44%

PROMISING

\$61k-\$351k

THEIR RANGE

8 REPOS

PRIVATE BUNDLE

2,546

COMMITTS

HUD'S OWN SCREEN - captured 31 AUG 2026

datavendor

Beta

Marketplace

Resources

Sign In

Sign Up

← New check

GitHub estimate

\$61,067 – \$350,557

Point estimate \$89,481

Useful technical assets, with gaps in coverage, documentation, or review history.

Each repository is scored on the same measurements HUD quality control grades: line coverage, pull request count, complexity, commits, documentation volume and reach, CI pass rate, and churn against complexity. The score grows the price exponentially from \$1,500 at zero to \$100,000 at the top of the scale. A measurement HUD could not take is imputed rather than zeroed, and the range spans what the price would be if every imputed term scored nothing against everything it could have scored.

CODEBASE SCORE

44%

Promising

Analysis overview

CODEBASES

8

SOURCE LOC

Unavailable

COMMITTS AT REFS

2,546

OPEN PULL REQUESTS

20

CLOSED WITHOUT MERGE

1

CONTRIBUTOR ENTRIES

22

STORED ARCHIVE SIZE

2477 MB

LANGUAGES

11 (Batchfile, Dockerfile, HTML, JavaScript)

TOTAL PULL REQUESTS

88

MERGED PULL REQUESTS

67

TOTAL ISSUES

1

PUBLIC DISCOVERABILITY

Not public

Totals are sums across the selected repository analyses. The same person can appear as a contributor in more than one repository. Source LOC is the exact count of non-blank lines in recognized programming-language files plus JSON, TOML, and YAML at the selected ref, including checked-in generated or vendored code. It is unavailable for public

Point \$89,481 | Score 44% Promising | Range \$61,067-\$350,557 | Discoverability: not public. HUD imputes missing measurements; \$350,557 is not cash I will cite. This screen transfers no rights. A payout exists only after a signed licence and money in the bank.

CONFIDENTIAL | HUD / DataVendor match only | Not a sale of the Lab

1 / 4

THE GRANT
Non-exclusive licence to a copy of the eight private repositories, for model-training and related evaluation by a matched lab through HUD / DataVendor. Same corpus may be licensed again. Each licence is a separate payment.

I RETAIN
Copyright. The Aldeazz AI Lab brand. Domains. The right to operate, improve, and sell every product. The running systems. Future commits after the licensed snapshot.

OUT OF SCOPE
Runtime and secrets (.env, keys, tokens). HubSpot CRM. Telegram and WhatsApp chat logs. Customer PII. Third-party or client material. Hostnames and internal ids. aideazz stays public: it deploys <https://aideazz.xyz/portfolio>

THE DEAL
A matched lab gets a copy. I keep the Lab. Lazarus / Turing already has this shape.

THE EIGHT REPOS - HUD'S POINT ESTIMATE

Repository	Score	Point	Range	Training signal
VibeJobHunterAIPA_AIMCF	58%	\$17,020	\$12.7k-\$50.6k	LangGraph hunter; 556 commits; 54k LOC; 131-test eval
EspaLuzWhatsApp	57%	\$16,735	\$12.4k-\$49.7k	WhatsApp advisor; 388 commits; bilingual tool-use
EspaLuz_Influencer	57%	\$16,662	\$12.9k-\$45.2k	Publishing agent; 126 commits; 100% CI
EspaLuzFamilybot	52%	\$13,303	\$10.8k-\$32.0k	Telegram advisor; 207 commits; 100% CI
AIPA_AITCF	47%	\$10,910	\$5.5k-\$64.5k	Ops brain; 1,052 commits; five-provider LLM chain
dragontrade-agent	35%	\$6,450	\$3.0k-\$47.2k	Human-on-last-click paper trading
AILA	25%	\$4,264	\$2.0k-\$31.2k	Personal AI (paused)
atlas-captures	24%	\$4,137	\$1.9k-\$30.2k	Ad-angle time-series data
EIGHT TOGETHER	44%	\$89,481	\$61.1k-\$350.6k	2,546 commits; 88 PRs (67 merged); 247.7 MB

Almost every repo has test coverage imputed. I do not ask you to treat imputed-perfect (\$350k) as an offer.

- LIVE PROOF (not in the licensed copy)
- <https://aideazz.xyz/portfolio> (primary hub)
 - <https://aideazz.xyz/api> (audit API + free tool)
 - <https://webhook.aideazz.xyz/whitespace/atlas.html> (Atlas board)
 - <https://aideazz.xyz/portfolio#portfolio-inquiry-form>
 - <https://aideazz.xyz/sop-ai-ops.html> (ops SOP)
 - <https://aideazz.xyz/blog>
 - <https://podcast.aideazz.xyz/>

WHAT I WILL NOT SIGN

Exclusive grant. Copyright assignment. Any term that requires taking a product down, transferring a domain, or including CRM / chat logs / keys / client data. An exclusive HUD paper would poison the Lazarus / Turing cheque.

ASK
Match a lab to this bundle. Tell me the actual payout. Licence is non-exclusive. The Lab stays up. I attend the DataVendor call. I sign a marketplace NDA that keeps ownership. I do not sign exclusive.

Elena Revicheva | Aldeazz AI Lab | Panama City | UTC-5
aipa@aideazz.xyz | <https://aideazz.xyz/portfolio>
Call: 31 Aug 2026, 09:40 America/Panama | cal.com/team/hud/talk-to-us-data-vendor-platform
This brief is not source code. The licensed copy is the private GitHub snapshot HUD already archived for the estimate.

WHY EACH SYSTEM EXISTS - AND HOW IT RUNS NOW

Provenance for a training lab, not a hiring dossier. Each of these eight started as a real constraint on one operator, one server. Live products stay up. This packet licences a copy of source.

Shared spine: one Oracle Cloud box, Telegram as the control room, HubSpot as the ledger, TypeScript where money/CRM state lives, Python where the ML stack lives, and a five-provider LLM chain so a zero-balance outage cannot stop the fleet. On 19 Aug Anthropic hit zero; the chain fell through to OpenAI in 2.4 seconds and nothing dropped.

EspaLuzFamilybot**Bilingual expat advisor - Telegram****WHY IT WAS BUILT**

Four people arrived in Panama at once without Spanish: a child, two retirees, one adult solving all of it. Language apps teach vocabulary on a schedule. The need was the sentence in front of you, plus which form, which queue, what you are actually agreeing to. Every arriving household has that first month.

HOW IT RUNS NOW

Live on Telegram. Per-family memory, not per-user: LangChain for the current conversation, pgvector RAG on Oracle Autonomous DB for everything durable ("which queue did I use last time?"). Voice in and out, image understanding for forms, payments wired for trials. Emotional intelligence before grammar. Designed for 21 Spanish-speaking countries, not Panama-only. Engineering is finished; families use it.

EspaLuzWhatsApp**The same advisor - WhatsApp****WHY IT WAS BUILT**

In Latin America WhatsApp is the internet. Telegram is where early adopters are. One advisor, two front ends, or the product only reaches half the household.

HOW IT RUNS NOW

Live on WhatsApp. Shared logic with the Telegram twin. Same memory, voice, and image path. Two clients over one brain so a family is one record, not two bots that forget each other.

EspaLuz_Influencer**Publishing agent - blog, social, cadence****WHY IT WAS BUILT**

One operator cannot also post to LinkedIn, Instagram, X and a blog every day. The Lab needed a public face without performing on social media. Only things that actually happened, or it is noise.

HOW IT RUNS NOW

Daily article at 14:30 Panama, cross-posted with canonical URLs. Bilingual social atoms on a schedule. Grounded generation: production evidence is collected before the model writes (uptimes, commits, CRM counts) and a number that cannot be traced is refused. Every link carries UTM and Atlas IDs.

VibeJobHunterAIPA_AIMCF**Scored discovery with a human on the last click****WHY IT WAS BUILT**

Mechanical search is volume. Judgment is not. The product problem is scored routing that a human can interrupt, plus a judge that can be audited instead of trusted.

HOW IT RUNS NOW

LangGraph state machine: gate, score, route, notify. SQLite checkpointing so a crash resumes. Auto-apply was removed after it produced volume and zero outcomes. Scoring is an LLM judge with two counterweights: a 131-test 4-layer eval harness (keyword, bias, golden-set, LLM-as-judge), and a feedback loop from real accept/reject reasons. Fails closed on an LLM outage. 11 sources, hourly.

AIPA_AITCF**Ops brain - sales operation + fleet control****WHY IT WAS BUILT**

No sales team, no Panama network, no budget, and eight systems that can break overnight. Customers and uptime both had to be software.

HOW IT RUNS NOW

Telegram is the control plane (82 commands, 08:00 briefing). HubSpot is the ledger. A webhook API is the front door. Form submission becomes a CRM record, a drafted reply, and a one-tap approval card. A synthetic lead hits the real pipeline daily at 07:45: silent on pass, alerting on failure. /code and /fix open real GitHub PRs; a human is on the last click. This repo also holds the Atuona automated film pipeline (clips in, mixed film out, nobody on a timeline) and the Atlas radar (detect, not predict).

dragontrade-agent**Human-on-last-click paper trading****WHY IT WAS BUILT**

An agent that can take an irreversible action needs a different design than a chatbot. Built to learn that constraint: trading education and paper trading only. Never live money.

HOW IT RUNS NOW

The agent proposes; a human disposes. Own dashboard. Paper trading only. Same last-click rule as the rest of the fleet.

atlas-captures**Ad-angle time-series data****WHY IT WAS BUILT**

Cold outreach with nothing to say fails. Public ad libraries are open data almost nobody reads systematically. A business will read a letter about their market, not a letter about the sender.

HOW IT RUNS NOW

Stored captures of which creative angle is observably running, and for how long. Detect, not predict: the system refuses to forecast. Weekly across seven service lanes. The live board is at <https://webhook.aideazz.xyz/whitespace/atlas.html>. This repo is the data HUD grades, not only the code.

AILA**Personal AI - paused****WHY IT WAS BUILT**

A life-scale assistant is a different product from the ops bots: governed longitudinal memory, safety, and continuity across months, not another Telegram command surface.

HOW IT RUNS NOW

Paused. Blueprint and inheritance notes exist (emotional/associative layers from Atuona; routing and persistence craft from AIPA). Thin in the HUD bundle because it is not the production control room. Included because it is a distinct private codebase.